

time_series

Contents

1	Read Time Series Date	1
2	Plotting Time Series	2
3	Decomposing Time Series	2
3.1	Decomposing Non-Seasonal Date	2
3.2	Decomposing Seasonal Date	3
3.3	Seasonally Adjusting	4
4	Forecasts using Exponential Smoothing	5
4.1	Simple Exponential Smoothing	5
4.2	Specify the initial value for the level in the HW model	7
4.3	forecast.HoltWinters()	7
5	Holt's Exponential Smoothing	11

1 Read Time Series Date

- scan() : read data into a vector or list from the console or file
- ts() : store the data into time series format

```
# alt+ctrl+i = create a new code chunk
# import data
kings <- scan("http://robjhyndman.com/tsdldata/misc/kings.dat",skip=3)
kings

## [1] 60 43 67 50 56 42 50 65 68 43 65 34 47 34 49 41 13 35 53 56 16 43 69 59 48
## [26] 59 86 55 68 51 33 49 67 77 81 67 71 81 68 70 77 56

# store the data in a time series object
# ts(), frequency = 'the time interval data collected', such as 12 for montly, 4 for quarter
# start(1982,2) = the data start from 1982 second quarter
kingstimeseries <- ts(kings)
kingstimeseries

## Time Series:
## Start = 1
## End = 42
## Frequency = 1
## [1] 60 43 67 50 56 42 50 65 68 43 65 34 47 34 49 41 13 35 53 56 16 43 69 59 48
## [26] 59 86 55 68 51 33 49 67 77 81 67 71 81 68 70 77 56


# the example of advanced use
births <- scan("http://robjhyndman.com/tsdldata/data/nybirths.dat")
births.ts <- ts(births, frequency = 12,start = c(1946,1))
knitr::kable(head(births.ts), 'markdown')
```

x
26.663
23.598
26.931
24.740
25.806
24.364

2 Plotting Time Series

- additive model : be able to describe seasonal data but trend
 - trend data can be transformed into log to decrease the trend effect.

```
plot.ts(kingstimeseries)
```




3 Decomposing Time Series

- to separate the data into its constituent components

3.1 Decomposing Non-Seasonal Date

- a non-seasonal time series consists of **a trend component** and **an irregular component**
- to estimate the trend component of a non-seasonal time series that can be described using an additive model : a smoothing method, such as MA.
- SMA() :
 - in “TTR” package
 - is used to smooth time series data, a simple moving average
 - specify the order of the simple moving average, n
 - * to estimate the trend component, might try smoothing the data with a simple moving average of a higher order
 - * trial-and-error

```
library(TTR)
kingstimeseriesSMA3 <- SMA(kingstimeseries, n=3)
plot.ts(kingstimeseriesSMA3)
```




3.2 Decomposing Seasonal Date

- a seasonal data consists of a **trend component**, a **seasonal component** and an **irregular component**.
- to estimate the trend component and seasonal component, using additive model
 - `decompose()` : decompose a time series into **seasonal**, **trend** and **irregular components**, using moving averages.
 - * return a list object

```
births.ts.decom <- decompose(births.ts)
plot(births.ts.decom)
```


Decomposition of additive time series



3.3 Seasonally Adjusting

- if the seasonal time series can be described by an additive model, the data can be seasonally adjusted by estimating the seasonal component and subtracting the estimated seasonal component from the original time series.
- the seasonally adjusted result just contains **the trend component** and ****an irregular component***
- `decompose()`

```
births.ts.decom.seasonallyadjusted <- births.ts - births.ts.decom$seasonal
plot(births.ts.decom.seasonallyadjusted)
```




4 Forecasts using Exponential Smoothing

- being used to make short-time forecasts for time series data

4.1 Simple Exponential Smoothing

- if the data can be described using an additive model with **constant level** and **no seasonality**
- using simple exponential smoothing to make short-time forecasts
- alpha : lies between 0 and 1
 - close to 0 : meaning little weight is placed on the most recent observations when making forecasts of future values
- HoltWinters() : simple exponential smoothing
 - beta = FALSE ; gamma = FALSE : for Holt's exponential smoothing or Holt-Winters exponential smoothing
 - **fitted** : the forecasts made by HoltWinters is stored in fitted

```
rain <- scan("http://robjhyndman.com/tsdldata/hurst/precip1.dat", skip=1)
rain.ts <- ts(rain, start = 1813)
plot.ts(rain.ts)
```




```
rain.ts.forecast <- HoltWinters(rain.ts,beta=FALSE,gamma=FALSE)
knitr::kable(head(rain.ts.forecast$fitted))
```

xhat	level
23.56000	23.56000
23.62054	23.62054
23.57808	23.57808
23.76290	23.76290
23.76017	23.76017
23.76306	23.76306

```
plot(rain.ts.forecast)
```

Holt–Winters filtering



- the accuracy offorcasting : the sum of squared errors
– $\$SSE$

4.2 Specify the inital value for the level in the HW model

- `l.start`

4.3 `forecast.HoltWinters()`


- make forecast for further time point
- *forecast* package
- `h` : how many further time points going to make forecast **`forecast.HoltWinter()` could not find in the package `stats::forecast.HoltWinters()` is not availibale for R 4.0 `forecast::forecast()` has the same function**
- the predicted area, 80% and 95% prediction interval

```
require('forecast')
rain.ts.forecast2 <- forecast::forecast(rain.ts.forecast, h=8)
knitr::kable(rain.ts.forecast2, 'markdown')
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
1913	24.67819	19.17493	30.18145	16.26169	33.09470
1914	24.67819	19.17333	30.18305	16.25924	33.09715
1915	24.67819	19.17173	30.18465	16.25679	33.09960
1916	24.67819	19.17013	30.18625	16.25434	33.10204
1917	24.67819	19.16853	30.18785	16.25190	33.10449
1918	24.67819	19.16694	30.18945	16.24945	33.10694
1919	24.67819	19.16534	30.19105	16.24701	33.10938
1920	24.67819	19.16374	30.19265	16.24456	33.11182

```
plot(rain.ts.forecast2)
```


Forecasts from HoltWinters



- if there are correlations between **forecast errors** for **successive predictions**, it is likely that the simple exponential smoothing forecasts could be improved upon by another forecasting technique.
- forecast error : a measure of accuracy of the predictive model is SSE, sum of squared error for the **in-sample** forecast errors. it is stored in **residual** of the forecast of HoltWinters.
- if the SSE could be improved, the predictive model should be improved as well.
- `acf()` : showing a correlogram of the **in-sample** forecast errors, to specify the max lag

```
acf(rain.ts.forecast2$residuals, lag.max = 20, na.action=na.pass)
```


Series rain.ts.forecast2\$residuals



```
#na.action=na.pass to pass the NA data
```


- the lag3 is just touching the significance bounds
- **Ljung-Box test** : to test whether there is significant evidence for non-zero correlations at lags 1-20.
- Box.test() : Ljung-Box test

```
Box.test(rain.ts.forecast2$residuals, lag = 20, type = 'Ljung-Box')
```

```
##  
## Box-Ljung test  
##  
## data: rain.ts.forecast2$residuals  
## X-squared = 17.401, df = 20, p-value = 0.6268
```

- the p-value is 0.6, so there is a little evidence of non-zero autocorrelation in the in-sample forecast error at lags 1-20.
- checking the normal distribution with mean zero and constant variance : to be sure that the predictive model cannot be improved upon.


```
plot.ts(rain.ts.forecast2$residuals)
```



histogram plotting : to check whether the forecast errors are normally distributed with mean zero.

```
plotForecastError <- function(forecasterrors){
  forecasterrors <- na.omit(forecasterrors)
  # make a histogram of the forecast errors:
  mybinsize <- IQR(forecasterrors)/4
  mysd <- sd(forecasterrors)
  mymin <- min(forecasterrors) - mysd*5
  mymax <- max(forecasterrors) + mysd*3
  # generate normally distributed data with mean 0 and standard deviation mysd
  mynorm <- rnorm(10000, mean=0, sd=mysd)
  mymin2 <- min(mynorm)
  mymax2 <- max(mynorm)
  if (mymin2 < mymin) { mymin <- mymin2 }
  if (mymax2 > mymax) { mymax <- mymax2 }
  # make a red histogram of the forecast errors, with the normally distributed data overlaid:
  mybins <- seq(mymin, mymax, mybinsize)
  hist(forecasterrors, col="red", freq=FALSE, breaks=mybins)
  # freq=FALSE ensures the area under the histogram = 1
  # generate normally distributed data with mean 0 and standard deviation mysd
  myhist <- hist(mynorm, plot=FALSE, breaks=mybins)
  # plot the normal curve as a blue line on top of the histogram of forecast errors:
  points(myhist$mids, myhist$density, type="l", col="blue", lwd=2)
}
plotForecastError(rain.ts.forecast2$residuals)
```

Histogram of forecasterrors



- the Ljung-Box : little evidence of non-zero autocorrelation
- the normal distribution with mean zero
- conclusion : the simple exponential smoothing method provides an adequate predictive model.
- the 80%-95% predictive intervals are valid

5 Holt's Exponential Smoothing

- the dataset with trend (increasing or decreasing) and no seasonality and can be described by an additive model. The short-time forecasts can be made by Holt's exponential smoothing.
- Holt's exponential smoothing estimate the level and slope at the current time point.
- Smoothing :
 - alpha, for the estimate of **the level** at the current time point.
 - beta, for the estimate of **the slope of the trend component** at the current point

```
skirts <- scan("http://robjhyndman.com/tsdldata/roberts/skirts.dat", skip=5)
skirts.ts <- ts(skirts, start = 1866)
plot.ts(skirts.ts)
```



- `HoltWinters(gamma=FALSE)` : Holt's exponential smoothing

```
skirts.ts.fc <- HoltWinters(skirts.ts, gamma=FALSE)
skirts.ts.fc
```


```
## Holt-Winters exponential smoothing with trend and without seasonal component.
##
## Call:
## HoltWinters(x = skirts.ts, gamma = FALSE)
##
## Smoothing parameters:
##   alpha: 0.8383481
##   beta : 1
##   gamma: FALSE
##
## Coefficients:
##           [,1]
## a 529.308585
## b  5.690464
```

```
skirts.ts.fc$SSE
```

```
## [1] 16954.18
```

```
plot(skirts.ts.fc)
```

Holt-Winters filtering



* the estimate value of alpha and beta are both high, which mean the current value of the level and of the slop of the trend component are based mostly upon very recent observations in the time series. * the SSE is high cause the level and the slop of the time series change a lot over time

```
skirts.ts.fc2 <- forecast::forecast(skirts.ts.fc,h=19,na.action=na.pass)
plot(skirts.ts.fc2)
```

Forecasts from HoltWinters

